

Linear Regression built from scratch

Linear regression is a supervised learning algorithm that has a deep root from statistics. It is one of the machine learning algorithms that look quite simple yet can be surprisingly useful. In this article, we present the basics of a linear regression learner and build one from scratch using Python.

1 What is regression?

Regression is one type of supervised machine learning.

Well that does not appear to help a lot. Let's start with one real-life example instead. When we build a bookshelf, we have to choose which plywood to use. From common sense, we know if our bookshelf is wide, we need thicker or stronger plywood, otherwise the bookshelf would bow or sag because of the weight of books. In contrast, if our bookshelf is narrow, perhaps thinner or weaker plywood would be fine. In that case, using thicker or stronger plywood would add unnecessary cost and thus is not an optimal choice.

Here we have two groups of data: "plywood" and "width of bookshelf". Our task is to answer this question: given our desired width of bookshelf, which plywood should we use? If we call "width of bookshelf" as x and "plywood" as y , then our task is "given x , what is y ?"

How do we answer this question? Admittedly there are many ways, but probably the go-to method for most people is simply looking around existing bookshelves. We just look around, find existing bookshelves, gather their width and which plywood did they use, and then determine the plywood for our new bookshelves. The underlying rationale is that we assume our new bookshelf is similar in nature to these existing ones. If it has similar width as some of the existing ones, maybe we can just go ahead and use a similar type or thickness of plywood as those ones.

In other words, to answer the question "given x , what is y ", we just gather a bunch of existing x and y . We try to "learn" something from these existing x and y . Based on what we learned, we then make an educated guess of y for our new x . What we described here is a brief picture of supervised learning, in which the existing x and y are called *training* data. That is part of the story where the word *machinelearning* comes from, in which *machine* (computer) is doing the *learning* instead of us.

In the world of supervised learning, if y is regarded to a label or a type, such as "pine plywood", "cedar plywood", "redwood plywood", and so on, then this

problem is called “classification”. In contrast, if y is regarded to a continuous number, such as the thickness of a wood piece like “0.32 inch”, “0.44 inch”, “0.58 inch”, and so on, then this problem is called “regression”, which is the topic of our discussion today.

We admit this is not a great regression example since dimensions in traditional woodworking are usually limited to multiples of 1/8 inch, but assume we have a high-tech workshop using CNC machines to do woodworking and we can achieve any random dimension like 0.572 inch, then this is truly a regression problem.

2 Math expression of a linear regression

A regression problem, like we said above, is to answer “given x , what is y ” by looking at some existing pairs of x and y . In this setting, x is usually called *attribute* or *feature*, while y is usually called *response* or *outcome*.

Assuming we have n existing pairs of x and y , which we can written as

$$X = [x_1 \quad x_2 \quad \dots \quad x_n] \quad (1)$$

$$Y = [y_1 \quad y_2 \quad \dots \quad y_n] \quad (2)$$

The task of regression is to find the relation from x to y . In machine learning, sometimes people call this as a *function approximation*, emphasizing we are trying to approximate a function whose input is x and output is y . The function from X to Y can be any type, but in linear regression, as the name implies, we assume this function is linear, which means Y is approximately linear to X .

$$y \approx wx + b \quad (3)$$

In this equation, w is usually called *coefficient* in statistics and *weight* is machine learning, while b is usually called *intercept* in in statistics and *bias* is machine learning. The predicted or modeled outcome $wx + b$ is often denoted as \hat{y} .

$$\hat{y} = wx + b \quad (4)$$

In the simplest case, when X only has one attribute, like out example, “the width of bookshelf”, then w is one single scalar. This is often called *simple linear regression (SLR)* in statistics.

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \dots \\ \hat{y}_n \end{bmatrix} = w \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix} + b \quad (5)$$

For most of the cases, X often has more than one attribute. In our bookshelf example, maybe not only the width but also the height affects the choice of plywood. For a lot of real-world problems, we have dozens or even hundreds of

attributes and we might not know which attributes actually affect the outcome. In these cases, X is no longer a vector but a matrix, with rows representing different data *points* or *instances* and columns representing different data attributes or features. This is often called *multiple linear regression (MLR)* in statistics.

$$X = \begin{bmatrix} x_{1,1} & x_{1,2} & \dots & x_{1,k} \\ x_{2,1} & x_{2,2} & \dots & x_{2,k} \\ \dots & \dots & \dots & \dots \\ x_{n,1} & x_{n,2} & \dots & x_{n,k} \end{bmatrix} \quad (6)$$

In this X matrix, $x_{i,j}$ represents the j th *attribute* of the i th data *instance*. For example, $x_{1,1}$ represents the width (1st attribute) of the 1st bookshelf (1st instance) and $x_{1,2}$ represents the height (2nd attribute) of the 1st bookshelf (1st instance), while $x_{2,1}$ represents width of the 2nd bookshelf and $x_{2,2}$ represents the height of the 2nd bookshelf.

Accordingly, w is no longer a scalar but a vector with a length of k , where k is the number of attributes per instance.

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \dots \\ \hat{y}_n \end{bmatrix} = \begin{bmatrix} x_{1,1} & x_{1,2} & \dots & x_{1,k} \\ x_{2,1} & x_{2,2} & \dots & x_{2,k} \\ \dots & \dots & \dots & \dots \\ x_{n,1} & x_{n,2} & \dots & x_{n,k} \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \\ \dots \\ w_k \end{bmatrix} + b \quad (7)$$

For simplicity, we can also incorporate bias b into this vector by adding a column of 1 as the first column of X and replacing b with w_0 .

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \dots \\ \hat{y}_n \end{bmatrix} = \begin{bmatrix} 1 & x_{1,1} & x_{1,2} & \dots & x_{1,k} \\ 1 & x_{2,1} & x_{2,2} & \dots & x_{2,k} \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_{n,1} & x_{n,2} & \dots & x_{n,k} \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ \dots \\ w_k \end{bmatrix} \quad (8)$$

This equation can then be written as matrix forms.

$$\hat{Y} = Xw \quad (9)$$

Now we have a math representation of the linear regression problem. The true outcome is $Y = [y_1, y_2, \dots, y_n]$ and the predicted outcome is $\hat{Y} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_n]$. Our task is “finding the proper $w = [w_0, w_1, w_2, \dots, w_k]$ so that \hat{Y} is as close to Y as possible”.

3 How far away is \hat{Y} from Y ?

To achieve our task, the first thing we need to do is to define what do we mean by “close”. In other words, we need some quantitative measure to represent how far away is \hat{Y} from Y .

To define this quantitative measure, let’s use three different data sets plotted in Figure 1 as examples. For simplicity, these data sets have the same x with




Figure 1: y : true outcome

only one data attribute. The horizontal axis plots x and the vertical axis plots y .

First we ask this question: what is the worst prediction \hat{Y} a reasonable person can make?

Admittedly, we can make a lot of arbitrarily bad predictions of \hat{Y} that is absurdly not close to Y at all. However, as *reasonable* people, we can probably all agree that there is a baseline for \hat{Y} which is \bar{y} , the mean value of Y . In our bookshelf example, this means for any new bookshelf, we will just use the average plywood thickness of existing bookshelves, regardless of the width of the new bookshelf. This choice might be wrong, but the point is that it is not absurdly wrong. Usually, we can use this “wrong” answer as a baseline for further improvement.

In math expression, guess everything as the mean of Y is basically let w_0 be the average of Y and let all other w_i be 0.

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \dots \\ \hat{y}_n \end{bmatrix} = \begin{bmatrix} 1 & x_{1,1} & x_{1,2} & \dots & x_{1,k} \\ 1 & x_{2,1} & x_{2,2} & \dots & x_{2,k} \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_{n,1} & x_{n,2} & \dots & x_{n,k} \end{bmatrix} \begin{bmatrix} \bar{y} \\ 0 \\ 0 \\ \dots \\ 0 \end{bmatrix} = \begin{bmatrix} \bar{y} \\ \bar{y} \\ \bar{y} \\ \dots \\ \bar{y} \end{bmatrix} \quad (10)$$

In this case, our true outcome is $Y = [y_1, y_2, \dots, y_n]$ and our predicted outcome is $\hat{Y} = [\bar{y}, \bar{y}, \dots, \bar{y}]$. In Figure 2, the “distance” between these two is visualized as vertical lines connecting the corresponding y_i and \bar{y} .

This distance can be defined quantitatively as the sum of square of the length of these vertical lines and is denoted as SS_T , which stands for *total sum of squares*. In some way, we can think of SS_T as a baseline of our model. A reasonable model should be at least works better than this. In our example, all three data sets happen to have similar SS_T of around 4200 to 4300.

$$SS_T = \sum_{i=1}^n (y_i - \bar{y})^2 \quad (11)$$

In our model, we have $[w_0, w_1, w_2, \dots, w_k]$ which are different from $[\bar{y}, 0, 0, \dots, 0]$ and we somehow believe our choice is better. For now, let’s assume we somehow get these w_i magically and the predicted \hat{y} are plotted in Figure 3.




Figure 2: SS_T : distance between y and \bar{y}




Figure 3: \hat{y} : predicted outcome from linear regression

Here comes our second question: how much better is our model compared to baseline? Our model predicts the outcome as $\hat{Y} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_n]$ instead of $[\bar{y}, \bar{y}, \dots, \bar{y}]$, so the “distance” between our model and the baseline is SS_R , which stands for *regression sum of squares* (also called *explained sum of squares*).

$$SS_R = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 \quad (12)$$

Figure 4 visualized the concept of SS_R , which is the sum of squares of the length of vertical lines. We can think of SS_R as how much better our model is compared with the baseline. In other words, SS_R tells us how helpful our model is. For example, for the leftmost data set, SS_R is 105 and \hat{y} is not significantly different from \bar{y} . This tells us this model is not super helpful as there is no significant difference between using \hat{y} and simply using \bar{y} . In contrast, for the rightmost data set, SS_R is 3900 which tells us that this model is pretty useful as using \hat{y} is significantly different (hopefully better) than simply using \bar{y} .

Nevertheless, no matter how much better our model is, we probably will not achieve perfection. Our third question is how far away is our model from absolutely perfection?

Our model predicts the outcome as $[\hat{y}_1, \hat{y}_2, \dots, \hat{y}_n]$ while the prefect answer is the true outcome $[y_1, y_2, \dots, y_n]$. The distance from our model to absolute perfectness is SS_E , which stands for *error sum of squares*, (also called *residual*




Figure 4: SS_R : distance between \hat{y} and \bar{y}

sum of squares).

$$SS_E = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (13)$$

Figure 5 visualized the concept of SS_R , which is the sum of squares of the length of vertical lines. Smaller SS_E is, the less “error”, the closer our model is to perfectness. For example, SS_E of the rightmost case is 235, which is the smallest among the three cases. Correspondingly, \hat{y} of the rightmost case is very close to its y visually.




Figure 5: SS_E : distance between y and \hat{y}

For these three example data sets, we can list the sum of squares in the following table:

Table 1: Sum of squares of the three example data sets

SS_T	4301.12	4166.18	4172.50
SS_R	105.21	1135.85	3900.81
SS_E	4198.95	2983.73	235.20
$SS_T - SS_E - SS_R$	-3.04	46.60	36.49
$R^2 = 1 - \frac{SS_E}{SS_T}$	2.4%	28.4%	94.4%

For most of linear regression, in theory, it can be approved that

$$SS_T = SS_R + SS_E \quad (14)$$

In some sense, SS_T is the distance from the baseline to perfection, SS_R is the distance from the baseline to our model, and SS_E is the distance from our model to perfection.

A variable called R-Squared (R^2) is then defined as

$$R^2 = 1 - \frac{SS_E}{SS_T} = \frac{SS_R}{SS_T} \quad (15)$$

R-Squared is bound between 0 and 1. The closer to 0, the closer our model is to baseline; the closer to 1, the closer our model is to perfectness. For our example, as listed in the table, R-Squared is 2.4%, 28.4%, and 94.4% respectively.




Figure 6: R^2 : how much is explained by the regression

The higher R-Squared is, the more the data is *explained* or *reflected* in the prediction. If the R-Squared is too small, such as the leftmost case, then it probably indicates that linear regression is not a suitable tool for this problem. The rightmost case is usually *too good to be true* for most of the complex problems in real-world. Most of the time, especially if the data involve human behavior like economics, R-Squared of around 30% is often considered to be good.

Back to our task, now we can define our task as to find w that achieves the largest R-Squared possible. Since SS_T is a constant value of problem only involves y and \bar{y} regardless the choice of our model, our task is then simply to minimize SS_E .

4 How do we minimize SS_E ?

As we explained above, SS_E indicates the distance between our model and perfection, and thus we would like to minimize it.

$$SS_E = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (16)$$

where \hat{y}_i is the predicted outcome from our model

$$\hat{y}_i = w_0 + x_{i,1}w_1 + x_{i,2}w_2 + \dots + x_{i,k}w_k = \sum_{j=0}^k x_{i,j}w_j \quad (17)$$

Therefore, SS_E is written as

$$SS_E = \sum_{i=1}^n (y_i - \sum_{j=0}^k x_{i,j} w_j)^2 \quad (18)$$

The problem is then: given $x_{i,j}$ and y_i , find w_j that minimize SS_E . This is one example of optimization problems, which can be solved by many methods. For linear regression, there is one particular effective method: *gradient descent*.

The basic idea of gradient descent is to find the derivative between SS_E and w_j , and gradually changing w_j to the desired direction. In other words, if increase of w_j would decrease SS_E , we will increase w_j by a little bit and see how it goes. In contrast, if increase of w_j would increase SS_E , we will decrease w_j by a little bit and see how it goes.

$$\begin{aligned} \frac{\partial SS_E}{\partial w_j} &= \frac{\partial}{\partial w_j} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \frac{\partial}{\partial w_j} \sum_{i=1}^n (y_i - \sum_{j=0}^k x_{i,j} w_j)^2 \\ &= \sum_{i=1}^n \left[-2(y_i - \sum_{j=0}^k x_{i,j} w_j) \frac{\partial}{\partial w_j} \sum_{j=0}^k x_{i,j} w_j \right] \\ &= \sum_{i=1}^n \left[-2(y_i - \sum_{j=0}^k x_{i,j} w_j) x_{i,j} \right] \\ &= -2 \sum_{i=1}^n [(y_i - \hat{y}_i) x_{i,j}] \end{aligned} \quad (19)$$

This can also be written as

$$\begin{aligned} \frac{\partial SS_E}{\partial w_j} &= -2 \sum_{i=1}^n [(y_i - \hat{y}_i) x_{i,j}] \\ &= -2 [(y_1 - \hat{y}_1) x_{1,j} + (y_2 - \hat{y}_2) x_{2,j} + \dots + (y_n - \hat{y}_n) x_{n,j}] \end{aligned} \quad (20)$$

In matrix expression, this is

$$\frac{\partial SS_E}{\partial w_j} = -2(Y - \hat{Y})X \quad (21)$$

We can verify this expression as

$$\begin{aligned}
& \begin{bmatrix} \frac{\partial SS_E}{\partial w_0} \\ \frac{\partial SS_E}{\partial w_1} \\ \frac{\partial SS_E}{\partial w_2} \\ \dots \\ \frac{\partial SS_E}{\partial w_k} \end{bmatrix} = -2 \begin{bmatrix} y_1 - \hat{y}_1 & y_2 - \hat{y}_2 & \dots & y_n - \hat{y}_n \end{bmatrix} \begin{bmatrix} 1 & x_{1,1} & x_{1,2} & \dots & x_{1,k} \\ 1 & x_{2,1} & x_{2,2} & \dots & x_{2,k} \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_{n,1} & x_{n,2} & \dots & x_{n,k} \end{bmatrix} \\
& = -2 \begin{bmatrix} (y_1 - \hat{y}_1) \times 1 + (y_2 - \hat{y}_2) \times 1 + \dots + (y_n - \hat{y}_n) \times 1 \\ (y_1 - \hat{y}_1)x_{1,1} + (y_2 - \hat{y}_2)x_{2,1} + \dots + (y_n - \hat{y}_n)x_{n,1} \\ (y_1 - \hat{y}_1)x_{1,2} + (y_2 - \hat{y}_2)x_{2,2} + \dots + (y_n - \hat{y}_n)x_{n,2} \\ \dots \\ (y_1 - \hat{y}_1)x_{1,k} + (y_2 - \hat{y}_2)x_{2,k} + \dots + (y_n - \hat{y}_n)x_{n,k} \end{bmatrix}
\end{aligned} \tag{22}$$

One we get $\frac{\partial SS_E}{\partial w_j}$, we then update w_j accordingly. If $\frac{\partial SS_E}{\partial w_j} > 0$, this indicates an increase of w_j would increase SS_E . Therefore we should decrease w_j as the following:

$$w_j = w_j - \alpha \frac{\partial SS_E}{\partial w_j} \tag{23}$$

where α is an arbitrary coefficient to represent how far we want w_j to move along the desired direction.

In contrast, if $\frac{\partial SS_E}{\partial w_j} < 0$, this indicates an increase of w_j would decrease SS_E . Therefore we should increase w_j as the following:

$$w_j = w_j - \alpha \frac{\partial SS_E}{\partial w_j} \tag{24}$$

Notice in this equation, since $\frac{\partial SS_E}{\partial w_j} < 0$, w_j minus $\frac{\partial SS_E}{\partial w_j}$ is in fact increasing w_j .

This is one step of gradient descent. To find the optimal w , we iterate over and over again, moving w_j along the desired direction bit by bit from iteration to iteration. Eventually, every w_j converges to its optimal value. To check if gradient descent is converged, we can monitor the value of SS_E . If SS_E does not decrease from iteration to iteration or if the decrease is too small, we can conclude the iterations are converged and there is no need to continue updating w .

5 Implement gradient descent in Python

To implement linear regression and gradient descent in Python, first we construct a class called *LinearRegression*. The variable *max.iteration* represents

the allowable number of iterations so that the iterations would not run forever. The variable *tol* represents the tolerance of meaningful decrease of SS_E from one iteration to the next. If the decrease is less than this tolerance, then we conclude the decrease is not considerable anymore and we stop the iteration.

```
class LinearRegression:
    def __init__(self, x, y, max_iteration=1000, tol=1e-3):
        self.max_iteration = max_iteration
        self.tol = tol

    def gradient_descent(self, x, y, lr=0.01):
        # Get number of instances (n) and number of attributes (d)
        n = x.shape[0]
        d = x.shape[1]

        # Initialize weights and bias with zeros
        w = np.zeros(d)

        # Initialize list to store MSE
        mse_log = []

        # Loop for gradient descent
        for i in range(self.max_iteration):
            y_pred = np.matmul(x, w)
            mse = np.mean((y - y_pred) ** 2)

            # Terminate if decrease of MSE is less than tol
            if len(mse_log) != 0 and 0 <= (mse_log[-1] - mse) / mse_log[-1] < self.tol:
                break
            elif len(mse_log) != 0 and mse_log[-1] < mse:
                break
            else:
                mse_log.append(mse)

            gradient_w = -2 * np.matmul((y - y_pred), x) / n
            w -= lr * gradient_w

        return w, mse_log
```

One difference between the previous equations and the code is that we use MS_E in the code rather than SS_E . In the code, MS_E is simply defined as $\frac{SS_E}{n}$. As n is a constant for given data, the derivative of MS_E to w is then

$$\frac{\partial MS_E}{\partial w_j} = -\frac{2}{n}(Y - \hat{Y})X \quad (25)$$

In the code, this is implemented as

```
gradient_w = -2 * np.matmul((y - y_pred), x) / n
```

Another note is that the coefficient α is often called *learning rate* in machine learning, which is the variable *lr* in the code. Choosing proper learning rate is

crucial for a successful linear regression. We will talk about how we can actually determine learning rate later in this post.

6 Visualization of gradient descent

Let's use the third data set in our example data to visualize the gradient descent. For now, suppose we somehow can magically determine the proper learning rate.

In this example, the learning rate α is set as 0.007 for w_0 and 0.0001 for w_2 . At the beginning, we do not know much about w , so one method to initialize is to set all w_j to be zero. In our example, this means we start with $\hat{y} = 0x + 0$. In this case, we are far away from the true y with a SS_E of 15960. So we calculated the gradient of w_j using

$$\frac{\partial MS_E}{\partial w_j} = -\frac{2}{n}(Y - \hat{Y})X \quad (26)$$

In this example,

$$\begin{aligned} \frac{\partial MS_E}{\partial w_0} &= -30.406 \\ \frac{\partial MS_E}{\partial w_1} &= -2037.650 \end{aligned} \quad (27)$$

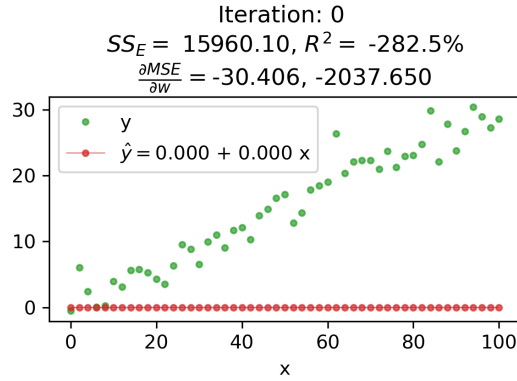


Figure 7: Iteration 0 of gradient descent

We need to update w_j to the direction that decrease MS_E .

$$\begin{aligned} w_0 &= w_0 - \alpha_0 \frac{\partial MS_E}{\partial w_0} = 0 - 0.007 \times (-30.406) = 0.213 \\ w_1 &= w_1 - \alpha_1 \frac{\partial MS_E}{\partial w_1} = 0 - 0.0001 \times (-2037.650) = 0.204 \end{aligned} \quad (28)$$

Now our model becomes $\hat{y} = 0.213 + 0.204x$. We move forward to the next iteration. With this new \hat{y} , our SS_E is reduced to 1807, indicating we are moving along the correct direction.




Figure 8: Iteration 1 of gradient descent

Again, we update w_j in this iteration.

$$\begin{aligned} w_0 &= w_0 - \alpha_0 \frac{\partial MSE}{\partial w_0} = 0.213 - 0.007 \times (-9.604) = 0.280 \\ w_1 &= w_1 - \alpha_1 \frac{\partial MSE}{\partial w_1} = 0.204 - 0.0001 \times (-644.348) = 0.268 \end{aligned} \tag{29}$$

After this iteration, our model becomes $\hat{y} = 0.280 + 0.268x$, our SS_E is reduced to 392, and we achieve a R-Squared over 90%. It feels pretty good. Let us move on to next iteration.




Figure 9: Iteration 2 of gradient descent

Again, we update w_j in this iteration.

$$\begin{aligned} w_0 &= w_0 - \alpha_0 \frac{\partial MSE}{\partial w_0} = 0.280 - 0.007 \times (-3.026) = 0.301 \\ w_1 &= w_1 - \alpha_1 \frac{\partial MSE}{\partial w_1} = 0.268 - 0.0001 \times (-203.764) = 0.289 \end{aligned} \quad (30)$$

Now our model becomes $\hat{y} = 0.301 + 0.289x$, our SS_E is reduced to 250, and we achieve a R-Squared of 94%. The fitted line of \hat{y} looks reasonable along the actual y . If we feel comfortable with this result, we can just stop here.

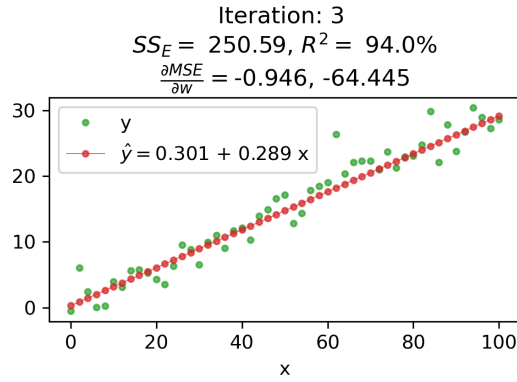


Figure 10: Iteration 3 of gradient descent

In this deliberately designed example, we achieved a pretty good linear regression model within just 3 iterations. Of course this is too good to be true for most real-world problems. To solve real-world problems, we have to examine how to choose a proper learning rate.

7 What is a proper learning rate?

In the previous section, we somehow magically get the learning rate for w_j . But we haven't talk about how to choose learning rate yet, so our next question is what is a proper learning rate? In gradient descent, we use learning rate α to update weights w . To successfully and effectively converge to an optimal solution, learning rate has to be proper. Figure 11 illustrates the effects of learning rates on one particular example data set. Different learning rates start from the same initial MSE but result in different learning behaviors.

If learning rate is too small, such as the curve corresponding to learning rate of 0.001, the learning process is unnecessarily slow and the iterations do not converge within the given allowable number of iterations. On the other hand, if learning rate is too large, then the learning process may become unstable as shown by the curve corresponding to learning rate of 0.2. The basic assumption




Figure 11: Learning curves of different learning rates

of partial derivative in calculus is that it is only applicable to a small region near that point, not applicable to regions far away from the point. Therefore, learning rate should be kept small enough so that the movement is within that region. In summary, learning rate needs to be a sweet spot, not too small nor too large, like the curve corresponding to learning rate of 0.1.

In our code, we use an automatic process to tune learning rate. We first start with a learning rate of 0.1, then we try 0.01, 0.001, and so on. We stop when the MSE achieved by a learning rate is more than 10 times larger than that of the previous learning rate. Then the learning rate corresponding to the minimum MSE is chosen as the optima learning rate.

```
def tune_learning_rate(self):
    lr = 0.1
    w, mse_log = self.gradient_descent(self.x_train, self.y_train, lr=lr)
    lr_list = [lr]
    mse_list = [mse_log[-1]]

    while lr > 1e-20:
        lr *= 0.1
        w, mse_log = self.gradient_descent(self.x_train, self.y_train, lr=lr)
        # if mse is 10 times larger than previous iteration, stop
        if mse_log[-1] > 10 * mse_list[-1]:
            break
        lr_list.append(lr)
        mse_list.append(mse_log[-1])

    opt_lr = lr_list[mse_list.index(min(mse_list))]

    return opt_lr
```

8 Scale data before training

For linear regression, it is not required to use scaled data. However, our code still scales and centers data before training.

The first reason is that this helps our autonomous tuning of learning rate. If the data are not scaled, the numerical range of different attributes may be significantly different and thus the proper learning rate for each attribute might be considerably different. By scaling the data, we are able to apply the same learning rate on all data attribute.

The second reason is that several advanced methods, which are based on linear regression, such as LASSO, ridge regression, and elastic net, all require the data to be scaled. To use our code as the foundation of these methods in the future, we apply data scaling and centering in our code.

```
class LinearRegression:
    def __init__(self, x, y, max_iteration=1000, tol=1e-3):
        self.max_iteration = max_iteration
        self.tol = tol

        self.x_raw = x
        self.mu = np.mean(x, axis=0)
        self.sigma = np.std(x, axis=0)
        self.x_train = self.standard_transform(x)
        self.y_train = y
        self.weights = None

    def standard_transform(self, x):
        scaled_x = (x - self.mu) / self.sigma
        scaled_x = np.insert(scaled_x, 0, np.ones(scaled_x.shape[0]), axis=1)
        return scaled_x
```

In this function, besides scaling and centering, we also insert a column of 1 as the first column of X so that we have a single w vector without using an additional variable of b for bias.

9 API of the learner

We use the following API to train and utilize the learner.

```
def fit(self, lr=None):
    if lr is None:
        lr = self.tune_learning_rate()
    w, mse_log = self.gradient_descent(self.x_train, self.y_train, lr=lr)
    self.weights = w

def predict(self, x):
    x_test = self.standard_transform(x)
    return np.matmul(x_test, self.weights)
```

Using this API, for given training and testing data sets, a learner can be built and trained using the following code:

```
learner = LinearRegression(x_train, y_train, plot=plot, verbose=True)
learner.fit()
```

To validate the learner by training data set, we can get predicted y of training x.

```
y_pred_train = learner.predict(x_train)
```

To utilize the learning on testing data set, we can get predicted y of testing x.

```
y_pred_test = learner.predict(x_test)
```

10 Example using *diamonds* data set

We use *diamonds* data set from *R* package *yarr* (<https://rdr.io/cran/yarr/man/diamonds.html>) as an example to run our learner. This data set has 150 data instances and 3 attributes: weight, clarity scale, and color scale. The outcome is the value of the diamond. The first 10 instances of the data set are previewed in the following table:

Table 2: Preview of *diamonds* data set

Weight	Clarity	Color	Value
9.35	0.88	4	182.5
11.10	1.05	5	191.2
8.65	0.85	6	175.7
10.43	1.15	5	195.2
10.62	0.92	5	181.6
12.35	0.44	4	182.9
11.24	1.09	6	195.9
9.77	1.43	4	188.6
11.83	0.95	6	190.3
9.55	1.05	5	190.7
...

Using our code, we train a linear regression model for this data set.

```
df = pd.read_csv('data/diamonds.csv', sep=' ')
df = df.to_numpy()
x = df[:, 0:-1]
y = df[:, -1]
learner = LinearRegression(x, y, plot=plot)
learner.fit()
y_pred = learner.predict(x)
```

The autonomous tuning of learning rates is shown in Figure 12. For this data set, learning rate of 0.1 seems to be the optimal choice as it achieved the minimum MSE and uses the smallest number of iterations.

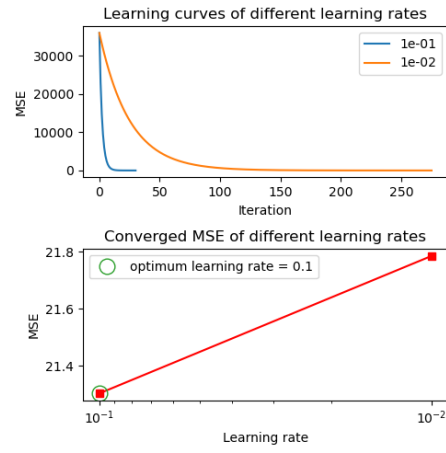


Figure 12: Tune of learning rate for *diamonds* data set

The results of the model can be visualized by plotting predicted values versus true values in Figure 13. The closer our predictions to the true values, the closer these plotted points to the red diagonal line. As we can see in the figure, the points in training and testing plots generally lay along the diagonal line.




Figure 13: Visualization of predicted values versus true values for *diamonds* data set

11 How do we interpolate the learning?

Now we have finished our linear regression on this *diamonds* data set. Imagine we need to give an elevator pitch to business partners or policy makers, what should we present? One of the most important outputs of linear regression is the weights of data attributes, which give us insights into the relations between attributes and outcome.

As we did gradient descent on scaled data, our weights also corresponds to the scaled data. To relate the weights back to original data, we need to do some calculation to derive the weights for unscaled data.

```
def get_unscaled_weights(self):
    unscaled_weights = self.weights[1:] / self.sigma
    unscaled_intercept = self.weights[0] - np.sum(self.weights[1:] * self.mu / self.sigma)
    unscaled_weights = np.insert(unscaled_weights, 0, unscaled_intercept)
    return unscaled_weights
```

Using this function, we obtain the weights for each attributes in the *diamonds* data set. As we explained before, the data outcome, which is the value of diamonds, can be approximated using a linear combination of weights and data attributes:

$$\text{diamond value} = 150.12 + 2.17(\text{weight}) + 20.67(\text{clarity}) - 0.63(\text{color}) \quad (31)$$

For this simple data set, we learn from the weights that the more the weights, the more the value, which seems reasonable. We also learn that the higher the clarity grade, the more the value, which also makes sense as people prefer diamonds with no or little inclusion or blemish. For color, the lower the color scale, the more the value, which probably indicates people prefer colorless diamonds.

12 How confident are we?

In our elevator pitch, we present the linear equation above and we talk about the relation between price and weight, clarity, or color. However, there is one an important question: how sure are we? In more technical words, in statistics, is our model significant enough?

To answer this question, we can perform a *significance test*. We still use n as the number of data instances and k as the number of data attributes. We already talked about SS_T , SS_R , and SS_E . One indicator is R-Squared which is already covered previously

$$R^2 = 1 - \frac{SS_E}{SS_T} = \frac{SS_R}{SS_T} \quad (32)$$

One drawback of using R-Squared is that it is compromised by increasing number of attributes. If the number of attributes is huge, then we can find w that generate a very small SS_E and thus a very large R-Squared. However, this does not always indicate that our model is perfect. As there are so many

random effects in these huge number of attributes, the model captured almost all of these random effect to achieve a high R-Squared. In machine learning, this is an example of *overfitting*. To overcome this, we introduce an adjusted version of R-Squared to deal with this.

To determine adjusted R-Square, we first define the corresponding mean of sum of squares.

$$MS_T = \frac{SS_T}{n-1} \quad (33)$$

$$MS_R = \frac{SS_T}{k} \quad (34)$$

$$MS_E = \frac{SS_E}{n-(k+1)} \quad (35)$$

Then the adjusted R-Squared is defined as

$$R_{\text{adjusted}}^2 = 1 - \frac{MS_E}{MS_T} \quad (36)$$

The relation between R-Squared and adjusted R-Squared can be derived.

$$\begin{aligned} R_{\text{adjusted}}^2 &= 1 - \frac{MS_E}{MS_T} \\ &= 1 - \frac{\frac{SS_E}{n-(k+1)}}{\frac{SS_T}{n-1}} \\ &= 1 - \frac{n-1}{n-(k+1)} \frac{SS_E}{SS_T} \\ &= 1 - \frac{n-1}{n-(k+1)} (1 - R^2) \end{aligned} \quad (37)$$

We can see that adjusted R-Squared decreases as k increases. In some sense, we can say using adjusted R-Squared encourages models with less data attributes. In other words, if two models achieve same R-Squared, but one with a significantly smaller number of attributes, then this one will have a smaller adjusted R-Squared and that tells us this model is better than the other one.

R-Squared and adjusted R-Squared can be calculated using the following code.

```
y = self.y_train
y_pred = np.matmul(self.x_train, self.weights)

n = self.x_train.shape[0]
k = self.x_train.shape[1] - 1

ss_total = np.sum((y - np.mean(y)) ** 2)
ms_total = ss_total / (n - 1)

ss_regression = np.sum((y_pred - np.mean(y)) ** 2)
```

```

ms_regression = ss_regression / k

ss_error = np.sum((y - y_pred) ** 2)
ms_error = ss_error / (n - k - 1)

r_square = 1 - ss_error / ss_total
adj_r_square = 1 - ms_error / ms_total

```

In our *diamonds* example, the R-Squared turns out to be 60.38% and the adjusted R-Squared turns out to be 59.28%.

Now we are ready to perform the significance test. For linear regression, we usually use a F-test to assess the significance. The *F-statistic* is defined as the ratio between the *explained* variances to *unexplained* variances.

$$F = \frac{MS_R}{MS_E} \quad (38)$$

Using this *F-statistic*, we can then determine the corresponding *P* value. To do that, first we need to determine the degree of freedom (dof) of the F-distribution. In linear regression, the first dof is k and the second dof is $n - (k + 1)$. In our *diamonds* example, we are looking at first dof of 3 and second dof of 108. The F-distribution and the corresponding survival function is plotted in the Figure 14.




Figure 14: F-test of linear regression on *diamonds* data set

When $F = 2.689$, the area under the probability density curve on the right side of 2.689 is 0.05, indicating a F value of 2.689 corresponds to a P value of 0.05. Similarly, the area under the probability density curve on the right side

of 2.135 is 0.10, indicating a corresponding P value of 0.10. In our *diamonds* example, our F turns out to be 54.883, which is way to the right of our plot. As a result, the P-value is extremely small.

P-value has strict definitions in statistics which we will not go into details today. What we need to know is that P-value indicates the probability that the reason why our model performs this well is purely by luck. In this case, this probability is very low, definitely smaller than 0.001. In statistics, 0.05 is usually used as the threshold. As a result, we conclude that this is not caused by pure luck, and thus, our model is at least somewhat “meaningful”, which is called “significant” in statistics.

13 Is every attribute relevant?

One final question we need to address is whether every attribute is relevant. By relevant, we mean “meaningfully contribute” to the outcome. If a certain attribute is random values not related to the outcome at all, we call that attribute irrelevant. For example, imagine we add a column of “weather of New York City when this diamond was produced” to the *diamonds* data set, this data attribute probably would not affect the value of diamonds at all. How do we identify if there is any irrelevant attribute?

Let us look back at our math expressions. This time we express everything as vectors and matrix. We are looking for a vector of weights w that linearly relates X to Y as close as possible. The “closeness” can be expressed as the sum square of error SS_E .

$$SS_E = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (39)$$

Written as matrix, this equation is

$$SS_E = (Y - \hat{Y})^T (Y - \hat{Y}) \quad (40)$$

The predicted outcome \hat{Y} of our model is obtained by

$$\hat{Y} = Xw \quad (41)$$

Therefore,

$$\begin{aligned} SS_E &= (Y - Xw)^T (Y - Xw) \\ &= (Y^T - w^T X^T)(Y - Xw) \\ &= Y^T Y - Y^T Xw - w^T X^T Y + w^T X^T Xw \end{aligned} \quad (42)$$

For the second term and the third term, we can find that they are transpose of each other.

$$(Y^T Xw)^T = w^T X^T Y \quad (43)$$

We then look at the dimensions of these items: Y^T is $1 \times n$, X is $n \times k$, and w is $k \times 1$. Therefore $Y^T X w$ is 1×1 . For a 1×1 matrix, its transpose is itself.

$$Y^T X w = w^T X^T Y \quad (44)$$

Therefore,

$$\begin{aligned} SS_E &= Y^T Y - Y^T X w - w^T X^T Y + w^T X^T X w \\ &= Y^T Y - 2w^T X^T Y + w^T X^T X w \end{aligned} \quad (45)$$

Following the principle of gradient descent, we look at the gradient between SS_E and w .

$$\begin{aligned} \nabla SS_E(w) &= \nabla Y^T Y - 2\nabla w^T X^T Y + \nabla w^T X^T X w \\ &= 0 - 2X^T Y + 2 \\ &= 2(X^T X w - X^T Y) \end{aligned} \quad (46)$$

To achieve the minimum SS_E , we need the gradient to be zero.

$$X^T X w - X^T Y = 0 \quad (47)$$

We then solve this equation for w . The w that makes the gradient to be zero is usually called \hat{w} .

$$X^T X \hat{w} = X^T Y \quad (48)$$

$$(X^T X)^{-1} X^T X \hat{w} = (X^T X)^{-1} X^T Y \quad (49)$$

$$\hat{w} = (X^T X)^{-1} X^T Y \quad (50)$$

In other words, among all possible w , \hat{w} is the one that makes $\hat{Y} = X \hat{w}$ as close to Y as possible. Put \hat{w} into the equation, we can get

$$\hat{Y} = X \hat{w} = X (X^T X)^{-1} X^T Y \quad (51)$$

If we let

$$H = X \hat{w} = X (X^T X)^{-1} X^T \quad (52)$$

then we have a transformation from Y to \hat{Y} .

$$\hat{Y} = [X (X^T X)^{-1} X^T] Y = H Y \quad (53)$$

Some people call H the *hat matrix* as it puts the hat on Y .

The term $(X^T X)^{-1}$ corresponds to the variances of w and is usually denoted as C .

$$C = MS_E (X^T X)^{-1} \quad (54)$$

In our *diamonds* example, the covariance matrix turns out to be

$$\begin{aligned}
C &= MS_E(X^T X)^{-1} \\
&= 21.87 \begin{bmatrix} 0.6020 & -0.0186 & -0.2190 & -0.0387 \\ -0.0186 & 0.0018 & -0.0014 & 0.0004 \\ -0.2190 & -0.0014 & 0.2103 & 0.0046 \\ -0.0387 & 0.0004 & 0.0046 & 0.0061 \end{bmatrix} \\
&= \begin{bmatrix} 13.164 & -0.408 & -4.789 & -0.846 \\ -0.408 & 0.040 & -0.031 & 0.008 \\ -4.789 & -0.031 & 4.599 & 0.101 \\ -0.846 & 0.008 & 0.101 & 0.133 \end{bmatrix}
\end{aligned} \tag{55}$$

The diagonal elements of C represent the standard deviation (se) of w . For each weight w_j , we have

$$se(w_j) = \sqrt{C_{j,j}} \tag{56}$$

Now we can perform a T-test on every w_j .

$$t_0 = \frac{w_j}{se(w_j)} \tag{57}$$

We already know our weights are

$$\begin{aligned}
w &= [w_0 \quad w_1 \quad w_2 \quad w_3] \\
&= [148.23 \quad 2.19 \quad 21.67 \quad -0.46]
\end{aligned} \tag{58}$$

Therefore,

$$\begin{aligned}
t_0 &= \begin{bmatrix} \frac{w_0}{se(w_0)} & \frac{w_1}{se(w_1)} & \frac{w_2}{se(w_2)} & \frac{w_3}{se(w_3)} \end{bmatrix} \\
&= \begin{bmatrix} \frac{w_0}{\sqrt{C_{0,0}}} & \frac{w_1}{\sqrt{C_{1,1}}} & \frac{w_2}{\sqrt{C_{2,2}}} & \frac{w_3}{\sqrt{C_{3,3}}} \end{bmatrix} \\
&= \begin{bmatrix} \frac{148.23}{\sqrt{13.164}} & \frac{2.19}{\sqrt{0.040}} & \frac{21.67}{\sqrt{4.599}} & \frac{-0.46}{\sqrt{0.133}} \end{bmatrix} \\
&= [40.85 \quad 10.93 \quad 10.10 \quad -1.26]
\end{aligned} \tag{59}$$

We then perform a t-test. As shown in the following figure, for a two-tailed t-test, a critical p-value of 0.05 corresponds to a t_0 value of 1.976 or -1.976.

Compare this to our four weights, we can determine that the first three weights have p-values significantly smaller than 0.05, indicating that it is very unlikely to get them by luck. In contrast, the last weight w_3 has a p-value of 0.21, which is greater than 0.05, and thus not significant in terms of statistics.

What does this tell us? The last weight, w_3 , corresponds to the data attribute color scale. A P-value larger than 0.05 indicates this attribute may not be relevant. In other words, we can interpolate this as “the value of a diamond is not affected by the color scale of that diamond”. Therefore, we probably should not include this attribute in our model. We will cover more advanced methods to properly deciding which attributes to include in our model in future posts.

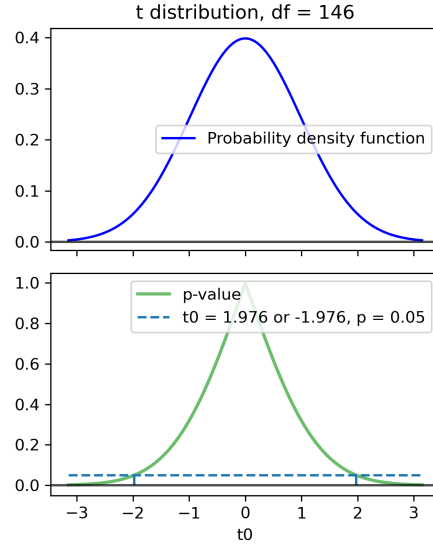


Figure 15: t-test of linear regression on *diamonds* data set

14 Compare our code to existing packages with more examples

Let's compare the results of our code to those of existing packages. We choose *sklearn.linear_model.LinearRegression* in *Python* and *lm* in *R* to compare.

14.1 diamonds

We already performed our liner regression code on *diamonds* and we got the following result with $R^2 = 63.67\%$.

$$\text{diamond value} = 148.23 + 2.19(\text{weight}) + 21.67(\text{clarity}) - 0.46(\text{color}) \quad (60)$$

Using *sklearn*, we get the following result with $R^2 = 63.72\%$.

$$\text{diamond value} = 148.34 + 2.19(\text{weight}) + 21.69(\text{clarity}) - 0.45(\text{color}) \quad (61)$$

Using *lm* in *R*, we get the following result with $R^2 = 63.73\%$.

$$\text{diamond value} = 148.33 + 2.19(\text{weight}) + 21.69(\text{clarity}) - 0.45(\text{color}) \quad (62)$$

As plotted in Figure 16, majority of these three results overlapp with each other.

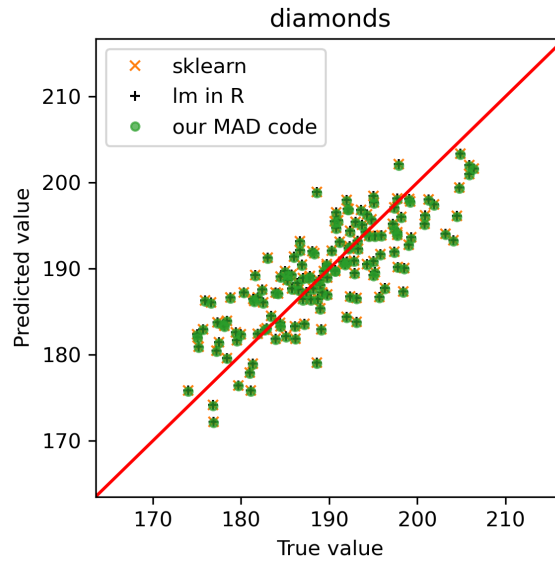


Figure 16: Compare our code to sklearn in Python and lm in R: *diamonds* data set

14.2 uscrime

The data set *uscrime* comes from <http://www.statsci.org/data/general/uscrime.html>. This data set has 47 instances with 15 attributes such as unemployment rate, median family income, and so on. The outcome is crime rate.

14.3 BostonHousing

The data set *BostonHousing* comes from <https://www.cs.toronto.edu/~delve/data/boston/bostonDetail.html>. This data set has 506 instances with 13 attributes such as crime rate, property tax rate, and so on. The outcome is median house value.

15 Summary

We talked about basics of linear regression, then built a linear regression learner which achieved similar results as those of existing packages.

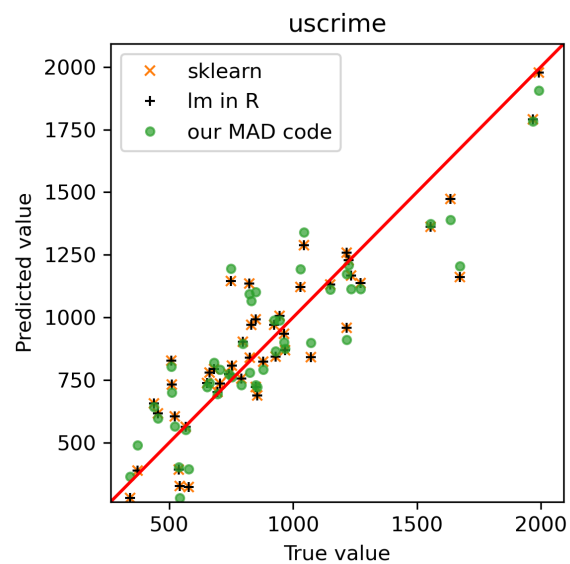


Figure 17: Compare our code to sklearn in Python and lm in R: *uscrime* data set

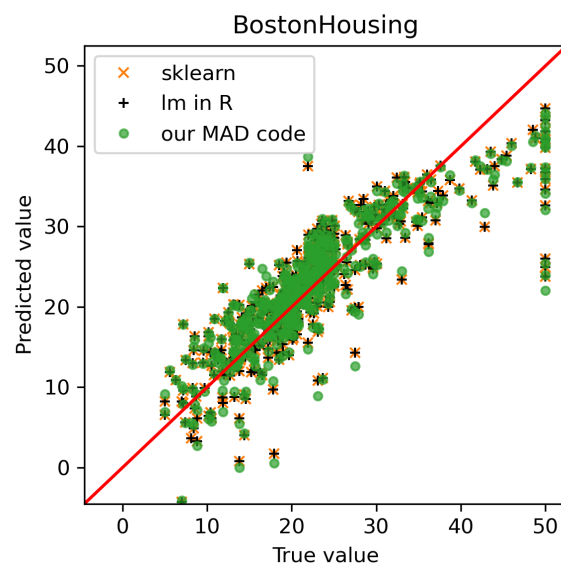


Figure 18: Compare our code to sklearn in Python and lm in R: *uscrime* data set